

Labyrinth Breach: A Reproducible Unity ML-Agents Benchmark for Asymmetric Multi-Agent Pursuit-Evasion in Dynamic Mazes

Ahmed Jawed, Imran Ashraf, Muhammad Sikander Raheem, Usman Irshad Bhatti, Arif Mahmood, Irene Delgado Noya, and Eduardo Garcia Villena

Abstract: We present Labyrinth Breach, a configurable Unity ML-Agents benchmark for asymmetric multi-agent pursuit-evasion in dynamic maze environments. Three Sentinel agents cooperatively pursue two Runner agents across a seeded training maze and procedurally generated held-out topologies with shifting walls, exit zones, and partial observability. Agents observe through vector states, a 360-degree planar horizontal ray fan, and last-known-position memory, and are trained with PPO and asymmetric self-play under a tactical reward-shaping framework designed to reduce free-rider collapse, cluster convergence, and passive evasion. Beyond the environment itself, we contribute an auditable evaluation pipeline that routes raw Unity logs into run-scoped artifacts, computes outcome and coordination metrics, and supports multi-seed seen/unseen layout evaluation and ablation studies. The empirical study includes five independently trained curriculum policy seeds, 30 canonical assist-off evaluation cells, 3,002 completed evaluation episodes, low-cost random and geometric-heuristic control policies, and two paired deployment interventions. The results show a balanced benchmark rather than a solved pursuit task: Runner wins remain common, held-out topology performance is comparable to seen-topology performance at the aggregate level, and dynamic walls alter tactical-event prevalence more clearly than terminal win rate. The source code and trained models are publicly available.¹ A demonstration video of trained agents in the dynamic maze is available online.²

Index Terms: Multi-agent reinforcement learning, pursuit-evasion, dynamic maze, asymmetric self-play, reward shaping, Unity ML-Agents, PPO

I. INTRODUCTION

Pursuit-evasion is a classical problem in robotics and multi-agent decision-making. One side attempts to capture while the other attempts to escape. The problem is a natural fit for

reinforcement learning (RL) because the desired behavior is difficult to hand-code when the environment involves adversarial interaction, obstacles, partial information, and multiple cooperating agents.

The original course project requires RL-based multi-agent pursuit-evasion using Unity ML-Agents [1]. We extend this requirement into a structured environment called *Labyrinth Breach*, where three Sentinels pursue two Runners inside a maze with exit zones and dynamic topology changes.

Simple chase environments do not capture the full difficulty of multi-agent coordination. In a richer maze setting, effective Sentinel behavior includes trapping, pincer movement, corridor denial, exit blocking, and route cutting. On the other side, effective Runner behavior includes distance management, path selection, survival, and escape through exits. These requirements align the project with robotics-style navigation and coordination problems rather than only direct-motion pursuit.

Our contributions are as follows:

- A configurable asymmetric 3v2 pursuit-evasion environment with a fixed training topology, seeded placements, dynamic walls, and procedurally generated held-out mazes.
- A tactical reward shaping framework (v5_active_agents) that targets free-rider collapse, clustering, and passive evasion through configurable shaping terms.
- A memory-augmented partial observation design combining a 360-degree planar horizontal ray fan with last-known-position tracking.
- A four-stage curriculum pipeline (fixed static → randomized static → low-frequency dynamic → high-frequency dynamic) with transfer between stages.
- A reproducible evaluation and evidence pipeline for multi-seed seen/unseen testing, coordination metrics, and ablations over memory, tactical rewards, dynamic walls, curriculum, and action assistance.

The study is organized around three research questions. **RQ1** asks whether a staged curriculum can produce competent policies as navigation complexity increases from an open arena to a dynamically changing maze. **RQ2** asks how strongly a fixed policy depends on maze geometry when transferred to held-out procedural layouts. **RQ3** asks which design elements, including memory, tactical reward shaping, wall dynamics, curriculum transfer, and action assistance, are responsible for observed performance and coordination. The

Ahmed Jawed, Muhammad Sikander Raheem, and Usman Irshad Bhatti are with the MS Artificial Intelligence program at Lahore University of Management Sciences (LUMS), Lahore, Pakistan. This work originated in AI641: AI for Robotics.

Imran Ashraf is with the Department of Information and Communication Engineering, Yeungnam University, Gyeongsan, Republic of Korea. Corresponding author: Imran Ashraf, imranashraf@yu.ac.kr.

Arif Mahmood is with the Department of Computer Science, Information Technology University, Lahore, Pakistan.

Irene Delgado Noya is with Universidad Europea del Atlantico, Santander, Spain, Universidade Internacional do Cuanza, Cuito, Angola, and Fundacion Universitaria Internacional de Colombia, Bogota, Colombia.

Eduardo Garcia Villena is with Universidad Europea del Atlantico, Santander, Spain, Universidad Internacional Iberoamericana Arecibo, Puerto Rico, USA, and Universidad de La Romana, La Romana, Dominican Republic.

¹<https://github.com/ahmedjawedaj/labyrinth-breach-marl>

²https://drive.google.com/file/d/1sxzktpvVNh_UYYxchMTZzzZNKLDyh09w/view

present study directly evaluates five independently trained curriculum policies on seen and held-out topologies and reports paired deployment interventions for action assistance and dynamic walls. Mechanism-level effects that require retraining, including memory removal, tactical-reward removal, and direct dynamic training, are registered as extensions and are not used to support the primary claims.

This scope discipline is important. The contribution is a reproducible simulation benchmark and an audited empirical study of asymmetric pursuit-evasion under dynamic topology. It is not presented as a new MARL algorithm, a hardware-robotics validation, or a state-of-the-art controller claim. This framing keeps each conclusion tied to an artifact, metric, and evaluation condition.

II. RELATED WORK

We organize the related work into five themes and compare the closest systems in Table I. The comparison exposes the intended niche of Labyrinth Breach: asymmetric teams, continuous control, partial observability, goal-directed evasion, dynamic topology, and run-level provenance in one configurable Unity benchmark. Table II then positions the current empirical evidence against recent journal systems. The comparison is deliberately conservative: published pursuit-evasion environments differ in dynamics, sensing, team size, terminal definitions, and evaluation budgets, so cross-paper numbers are evidence of scope and difficulty rather than a controlled leaderboard. The comparison also makes clear what the present system does not yet provide, most importantly physical-robot validation.

A. Pursuit-Evasion as a Multi-Agent Problem

Direct pursuit-evasion papers share the same capture-versus-escape structure as our task. In harder settings, useful behaviors go beyond direct motion and include route blocking, coordinated pressure, spatial entrapment, and strategic escape timing [5], [8], [13], [16], [20], [27]. De Souza Jr. et al. train decentralized pursuit policies for non-holonomic agents and transfer them to three physical drones [28]. Gonultas and Isler incorporate vehicle dynamics and sensor constraints and deploy learned policies on ground robots [8]. ViPER uses graph attention for visibility-based search and includes an aerial hardware demonstration [13]. These systems set a stronger robotics validation bar than simulation alone.

Labyrinth Breach studies a complementary configuration: three cooperative pursuers against two learned, goal-directed evaders in a maze whose topology can change during an episode. Capture is not the sole objective because Runners can actively reach exits, and the 3v2 asymmetry creates both within-team coordination and between-team adaptation. The contribution is therefore not a claim to replace dynamics-faithful or hardware-validated pursuit systems. It is a configurable testbed for studying how environment structure, sensing, memory, and reward design interact under reproducible MARL evaluation.

Recent journal work also shows how pursuit-evasion has moved beyond simple open-field capture. Du et al. model

unauthorized-UAV pursuit in urban airspace with MARL [21]. Peng et al. add limited visual fields, building occlusion, graph attention, and normalizing-flow policies for multi-UAV capture [22]. Qu et al. use centralized training, decentralized execution, LSTM memory, multi-agent soft actor-critic, and multi-stage reward guidance for cooperative USV pursuit [23]. Xu and Dang analyze emergent pursuit behaviors in a bounded grid world and report a 99.9% pursuer success rate over 1,000 randomized trials, but under a simpler discrete, complete-information setting [24]. These papers strengthen the need to report environment assumptions alongside performance.

B. Cooperative and Mixed MARL

Multi-agent reinforcement learning faces a non-stationarity challenge because all agents update simultaneously. MADDPG addressed this through centralized training with decentralized execution [2]. Yu et al. subsequently showed that PPO-based methods can be strong cooperative baselines when implementation details and hyperparameters are controlled carefully [4]. Our two policy groups use parameter sharing within roles: all Sentinels share one policy, and all Runners share another. This exploits role homogeneity without forcing the two sides to share observations, objectives, or network weights.

Unity ML-Agents PPO is deliberately retained as the primary baseline because it is integrated with the simulator and straightforward to reproduce. This choice should not be interpreted as an algorithmic claim. MAPPO offers a centralized critic, and MA-POCA is designed for cooperative groups whose membership or participation may change. A journal-grade algorithm comparison should hold environment steps, network capacity, evaluation budgets, and seed counts constant. The present extended manuscript specifies that comparison as future empirical work rather than implying superiority from one PPO training trace.

C. Credit Assignment

When multiple agents cooperate under team-level objectives, one agent may contribute by blocking a route or creating pressure while another gets the final capture. Plain shared reward can be too coarse for learning such indirect contributions [3], [6], [7], [17]. Recent methods decouple teammate reward streams [31] or reason over individual, joint, and correlated agent subsets [33]. A 2026 model-based pursuit-evasion system further combines adversary uncertainty with mutual policy advising during imagined rollouts [34]. Labyrinth Breach takes a simpler environment-level approach: event detectors identify pincer pressure, corridor control, exit denial, enclosure, and capture, after which role-specific reward policies assign configured bonuses.

This mechanism makes the source of each shaping reward observable in `reward_audit.csv`, but it does not solve credit assignment in the formal counterfactual sense. A detector may reward correlated behavior that was not causally responsible for a capture. For that reason, the final ablation protocol removes the tactical layer entirely and reports both outcome effects and event-frequency changes. If the coordination metrics do not degrade when the layer is removed,

the claimed mechanism is unsupported even if total return changes.

D. Partial Observability

Pursuit-evasion becomes more meaningful when sensing is limited. With full global state, the task is easier and less realistic. Works on visibility-constrained pursuit-evasion show that line-of-sight, uncertainty, and memory fundamentally change the problem [8], [9], [13]. R2PS goes further by maintaining a belief over possible evader nodes and evaluating two pursuers against an optimal asynchronous evader on 10 unseen real-world graphs [15]. Its 300-graph training set and 500 tests per graph establish a stronger partial-observability standard than a last-known-position feature. Our observation pipeline uses vector states combined with ray-based visibility and last-known-position memory rather than global state or camera images. The memory is intentionally low-capacity: it stores the last observed relative position of a target and a recency signal. It cannot infer a hidden trajectory or represent multimodal beliefs, so recurrent or explicit-belief memory remains an important baseline.

E. Curricula, Benchmarks, and Evaluation Practice

Curriculum learning is especially relevant when sparse terminal rewards and adversarial co-adaptation make direct training unstable. Grupen et al. combine environment and behavioral curricula in pursuit-evasion and analyze emergent implicit signaling [29]. DualCL combines an intrinsic-parameter curriculum with an environment generator and evaluates three seeds with 3,000 episodes per reported score, including two out-of-distribution scenes [11]. MatrixWorld connects pursuit-evasion, safety-constrained action execution, co-evolution, and autotutorials in a lightweight grid benchmark [9]. Labyrinth Breach instead uses a manually specified progression that preserves network dimensions across stages, enabling checkpoint transfer while introducing walls and then topology shifts. The direct-dynamic ablation is therefore necessary to separate curriculum benefit from training budget.

Cross-topology claims require more than changing placement seeds. Equilibrium Policy Generalization trains graph policies over 152 graphs and evaluates 10 real-world plus 10 held-out Dungeon graphs with 500 trials per graph [14]. It also uses an equilibrium oracle, so its game-theoretic interpretation does not transfer to PPO self-play here. We consequently use the narrower term “held-out topology transfer” and report each topology separately. Zheng et al. show a complementary mechanism: obstacle-aware rewards can improve encirclement, but their five-seed, 200,000-episode study evaluates success, efficiency, team size, obstacle count, and component ablations together [12]. This motivates testing whether our trap reward changes terminal outcomes and capture efficiency rather than merely increasing shaped return.

Evaluation methodology is a separate concern from environment design. Gorsane et al. identify inconsistent seed counts, evaluation budgets, and uncertainty reporting across cooperative MARL studies and propose a standardized protocol [30]. Our official protocol follows the same principles

at smaller project scale: independent training seeds, fixed evaluation manifests, held-out tasks, per-seed reporting, confidence intervals, and preservation of raw artifacts. The validated diagnostic uses only one checkpoint pair and therefore cannot substitute for this protocol.

F. Hybrid and Assisted Control

Hybrid controllers combine learned actions with a geometric, model-based, or heuristic control signal. Akinmolayan et al. blend a classical interception vector with PPO and anneal the blend during training [32]. Labyrinth Breach contains a related optional controller feature: the Sentinel action can be blended with a direction toward the tracked Runner. The active diagnostic configuration uses a Sentinel pursuit strength of 0.15. This makes action assistance a potential confound rather than a minor implementation detail. The coefficient is disclosed and assist-off performance is reported.

G. Explainability and Policy Interpretation

Explainable reinforcement learning is increasingly relevant for autonomous robots because strong performance alone does not show why a learned policy acts as it does. Cetin et al. apply SHAP to deep RL counter-drone decisions and show how feature attribution can make an interception policy more transparent [25]. A recent XMARL survey frames interpretability at multiple levels: individual decisions, interaction mechanisms, team strategy emergence, and system-level performance [26]. Labyrinth Breach follows this multi-level framing with artifact-level explanations rather than image saliency. Raw observations and actions support individual inspection, replay events expose interaction mechanisms, coordination metrics summarize emergent team strategies, and paired ablations test system-level dependencies. Full SHAP analysis is left as an extension because our deployed PPO policies use role-specific continuous actions, ray groups, memory features, and multi-agent coupling. Applying SHAP without careful feature grouping would risk producing visually attractive but weak explanations.

III. PROBLEM FORMULATION

We model the environment as a Partially Observable Stochastic Game (POSG) defined by the tuple

$$\mathcal{G} = \langle \mathcal{I}, \mathcal{S}, \{\mathcal{A}_i\}, P, \{\Omega_i\}, O, \{R_i\}, \gamma \rangle, \quad (1)$$

where $\mathcal{I} = \mathcal{I}_S \cup \mathcal{I}_R$ contains three Sentinels and two Runners. The global state $s_t \in \mathcal{S}$ contains agent poses and velocities, capture/escape states, exits, procedural maze state, and the active wall set. Each action $a_{i,t} = [u_{i,t}^x, u_{i,t}^z] \in [-1, 1]^2$ controls planar movement. Agent i receives only $o_{i,t} = O_i(s_t) \in \Omega_i$, with 117 dimensions for a Sentinel and 129 for a Runner. The Unity transition kernel $P(s_{t+1} | s_t, \mathbf{a}_t)$ combines deterministic control code, collision handling, and seeded procedural events.

Agents are homogeneous within a role but not across roles. Parameter sharing therefore gives

$$\pi_{\theta_S}(a_{i,t} | o_{i,t}), \quad i \in \mathcal{I}_S, \quad \pi_{\theta_R}(a_{j,t} | o_{j,t}), \quad j \in \mathcal{I}_R. \quad (2)$$

TABLE I
COMPARISON WITH RECENT LEARNING-BASED PURSUIT-EVASION SYSTEMS AND BENCHMARKS. “DYNAMIC TOPOLOGY” MEANS OBSTACLES CHANGE CONNECTIVITY DURING AN EPISODE, NOT ONLY THAT AGENTS MOVE.

Work	Teams	Control	Sensing	Environment	Validation	Primary emphasis
De Souza Jr. et al. [28]	Multi-pursuer, one evader	Non-holonomic	Local ordered neighbors	Open pursuit area	Simulation + drones	Decentralized pursuit and formation
Kouzeghar et al. [10]	Pursuer + scout roles	Aerial	Local neighbors	Random obstacles	Simulation + drones	Pursuit and scouting
DualCL [11]	Multi-pursuer, one evader	3-D aerial	Global state	Generated obstacles	Simulation + drones	Dual curriculum and transfer
Gonultas and Isler [8]	1v1	Car-like dynamics	Sensor constrained	Indoor obstacles	Simulation + ground robots	Dynamics and sensing constraints
MatrixWorld [9]	Variable	Discrete grid	Configurable	Static grid topology	Simulation	Safety, co-evolution, autocurricula
ViPER [13]	Multi-pursuer, evaders	Planar/aerial	Visibility limited	Multiply connected maps	Simulation + aerial hardware	Distributed visibility clearing
EPG/R2PS [14], [15]	Graph pursuit	Discrete graph	Full/partial	Cross-graph	Simulation	Topology transfer and robust belief
Akinmolayan et al. [32]	4v3 extension	Hybrid continuous	Radar + communication	Bounded defense arena	Simulation	Classical/PPO blending and self-play
Labyrinth Breach	3v2 learned teams	Continuous planar	Rays + vector memory	Changing 3-D maze	Unity simulation	Topology shifts, exit evasion, auditable metrics

TABLE II
STATE-OF-THE-ART POSITIONING. REPORTED PERFORMANCE IS TAKEN FROM THE ORIGINAL PAPERS WHEN A DIRECTLY STATED NUMBER IS AVAILABLE. DIFFERENCES IN ENVIRONMENT, SENSING, DYNAMICS, AND EPISODE RULES PREVENT DIRECT RANKING.

Work	Setting	Reported evidence in source work	Position relative to Labyrinth Breach
Du et al. [21]	Urban unauthorized-UAV pursuit with MARL	Journal study of cooperative UAV pursuit in urban airspace	Stronger urban-airspace framing. Labyrinth adds learned 3v2 evaders, exit objectives, changing maze topology, and run-provenance audit
De Souza Jr. et al. [28]	Decentralized multi-drone pursuit	Simulation plus physical drone transfer for non-holonomic agents	Stronger sim-to-real evidence. Labyrinth studies asymmetric two-team evasion and dynamic maze connectivity
Peng et al. [22]	Multi-UAV pursuit with limited visual field and building occlusion	NAGC compared with MARL baselines in high-precision simulation, including ablations and generalization scenarios	Stronger algorithmic comparison. Labyrinth contributes a reproducible Unity benchmark with exit-directed evasion and artifact-gated evaluation
Qu et al. [23]	Cooperative USV pursuit using MASAC, CTDE, LSTM, and staged reward guidance	Simulation evidence for cooperative capture against a dynamic target	Stronger maritime-domain specificity. Labyrinth contributes multi-role adversarial learning and dynamic indoor-like topology
Xu and Dang [24]	Bounded 2-D grid pursuit-evasion with MARL behavior mining	99.9% pursuer success in 1,000 randomized trials plus clustered behavior analysis	Stronger quantitative behavior mining in a simpler complete-information grid. Labyrinth uses continuous control, partial observability, and two goal-directed evaders
Labyrinth Breach	Unity 3v2 dynamic maze, assist-off canonical evaluation	30 policy-layout cells and 3,002 episodes: 42.2% Sentinel win rate, 53.5% escape rate, 20.3% held-out pincer episode rate	Not a SOTA-performance claim. Evidence supports a difficult, balanced, auditable benchmark and motivates registered retraining extensions

For role $q \in \{S, R\}$, PPO seeks parameters maximizing the expected discounted return against an opponent-policy mixture μ_{-q} ,

$$J_q(\theta_q, \mu_{-q}) = \mathbb{E}_{\pi_{\theta_q}, \pi_{-q} \sim \mu_{-q}, P} \left[\sum_{t=0}^T \gamma^t r_{q,t} \right]. \quad (3)$$

Our self-play controller uses $\mu_{-q} = 0.5\delta_{\theta_{-q}^{\text{latest}}} + 0.5 \text{Unif}(\mathcal{H}_{-q})$, where $\mathcal{H}_{-q}$ is a window of the ten most recent opponent snapshots. One role learns at a time. The learning role changes every 100,000 behavior steps, snapshots are saved every 20,000 steps, and the deployed opponent is resampled every 2,000 steps. This limits direct latest-policy cycling

but does not make the game stationary. The environment is mixed-sum rather than strictly zero-sum. Terminal rewards are opposing, but role-specific shaping rewards need not sum to zero. Consequently, neither a Nash-equilibrium claim nor convergence to a game-theoretic solution follows from stable training curves or ELO values.

The Sentinel team wins when both Runners are captured before timeout. The Runner side wins when at least one active Runner reaches an exit zone or when the episode reaches the active rule-config timeout. Official dynamic-maze runs use 180 seconds. Excluded legacy plots may use the earlier 120-second regime. Agent components use environment-managed

termination rather than a serialized per-agent `MaxStep`. This detail matters because a trainer step is not equivalent to an episode step when five agents share two policies.

Let $c_j(t) \in \{0, 1\}$ indicate that Runner j is captured and $e_j(t) \in \{0, 1\}$ indicate escape. The terminal outcome is

$$y = \begin{cases} S, & \prod_{j \in \mathcal{I}_R} c_j(T) = 1, \\ R, & \max_j e_j(T) = 1 \text{ or } T = T_{\max}. \end{cases} \quad (4)$$

The episode-level Sentinel win estimator over n completed episodes is $\hat{p}_S = n^{-1} \sum_{k=1}^n \mathbb{I}[y_k = S]$. We use Wilson score intervals for this binary proportion because the diagnostic contains small cells and rates near zero or one, where a normal approximation is unreliable.

Dynamic topology is represented by a time-indexed traversability graph $G_t = (V, E_t)$. At a wall event, a seeded controller changes a bounded subset ΔE_t subject to a safety buffer around agents. Exit reachability follows from the generator invariant that spanning-tree passages are never converted into dynamic edges. It does not rely on the currently inert `ExitGuard` tag check. Thus $E_{t+1} = E_t \triangle \Delta E_t$. The policy does not observe G_t globally. It must detect local changes through rays and context features. This produces a different problem from static procedural randomization, where a layout changes only between episodes.

This formulation supports the study of asymmetric team interaction, coordinated pursuit, tactical escape behavior, partial observability, dynamic environment changes, and generalization to unseen layouts.

IV. SYSTEM DESIGN

Figure 1 summarizes the complete workflow. The core design choice is that simulation, learning, evaluation, and reporting are not treated as separate manual steps. Instead, every paper result is generated from run-scoped artifacts and auditable metadata. This is intended to reduce the common MARL failure mode in which training curves, screenshots, and hand-selected demonstrations cannot be traced back to reproducible seeds and configurations.

A. Environment Structure

The implementation includes four scenes of increasing complexity:

- **Open arena baseline:** flat arena with no walls.
- **Static maze:** a fixed corridor topology with configurable placements.
- **Dynamic maze:** walls shift at configurable intervals during episodes.
- **Unseen evaluation maze:** novel layouts not encountered during training.

The dynamic maze is the core environment. Training uses a fixed base topology, while held-out evaluation topologies are generated deterministically from reserved seeds. In both cases, selected wall pillars can rise or lower on a timer. These shifts invalidate a fixed route during an episode and create trapping opportunities and escape adaptations in real time. Exit zones are placed explicitly for the Runner side, making evasion a goal-directed navigation task rather than pure survival.

Figure 2 provides the environment visual requested by reviewers. It emphasizes the state variables and behavioral events used by the evaluation pipeline rather than presenting a decorative screenshot. This makes the maze, agent roles, dynamic-wall mechanism, and qualitative behavior categories clear before the training details are introduced.

B. Seeded Topology Generation

Held-out mazes use a 13×13 occupancy grid containing a 6×6 lattice of candidate free cells. A seeded randomized depth-first search carves a spanning tree E_{tree} over all 36 cells. Four additional non-tree connections are opened permanently to reduce single-corridor degeneracy, and up to five further non-tree connections are marked as dynamic wall edges. The traversable graph is therefore

$$E_t = E_{\text{tree}} \cup E_{\text{loop}} \cup E_{\text{dyn}}(t), \quad E_{\text{dyn}}(t) \subseteq E_{\text{optional}}. \quad (5)$$

Because E_{tree} is never removed, every generated topology remains connected even when all dynamic connections are raised. This is a construction-level safety property, not an empirical assumption. The reserved topology seeds are 101, 202, 303, 404, and 505. Topology seed q_l is fixed within policy-layout cell l , whereas placement seed $q_l + k$ varies with episode k . Each episode row records both `maze_layout_id` and `maze_topology_seed`. This prevents placement randomization from being mislabeled as topology diversity.

C. Agent Architecture

The codebase separates agent logic, environment control, reward computation, sensing, and logging into distinct modules. `BaseAgent` handles movement, observation collection, and ML-Agents integration. `SentinelAgent` extends it with capture-radius logic and target tracking. `RunnerAgent` extends it with escape and captured-state management. The central `PursuitEvasionEnvController` manages episode lifecycle, capture detection, reward triggering, and terminal conditions.

Figure 3 makes the model boundary explicit. The neural policy consumes grouped vector observations and outputs a two-dimensional continuous movement command, while Unity handles collision, capture, escape, optional assistance, and artifact logging outside the network. This separation is important because the optional geometric assistance and the evaluation-only baselines are controller-level interventions, not additional learned layers.

The controller can optionally blend a learned direction $a_{i,t}$ with a geometric direction $h_{i,t}$:

$$\tilde{a}_{i,t} = (1 - \alpha_{i,t})a_{i,t} + \alpha_{i,t}h_{i,t}, \quad (6)$$

$$\alpha_{i,t} = \text{clip}(\rho_q[c_q + (1 - c_q)u_{i,t}], 0, 1). \quad (7)$$

where q is the role, ρ_q is the configured assist strength, and $u_{i,t}$ increases as the visible or remembered target becomes nearer. The source uses $c_S = 0.35$ for pursuit and $c_R = 0.30$ for evasion, with obstacle-aware sliding applied to $h_{i,t}$. Training uses $\rho_S = 0.15$ and $\rho_R = 0$, but the canonical publication evaluation sets both to zero so reported policy behavior is




Fig. 1. Proposed methodology and evidence flow. The same seeded configuration controls training, evaluation, logging, KPI extraction, statistical analysis, and manuscript figures. Strict artifact validation excludes runs with missing logs, unknown seeds, or mismatched scene/rule/reward provenance from publication tables.




Fig. 2. Environment and qualitative behavior schematic. The panels show the core visual elements that matter for the method: dynamic walls, exit zones, local sensing, pincer pressure, exit denial, and route change after topology shifts. The figure is schematic rather than a raw Unity screenshot so the experimental variables remain legible in print.




Fig. 3. Role-specific PPO architecture used by the Sentinel and Runner policies. The figure describes the implemented ML-Agents actor-critic network and runtime controller path. It is not a transformer or attention architecture.

unassisted. Since $h_{i,t}$ is computed outside the policy network, the matched assist-on intervention reinstates $\rho_S = 0.15$ and measures a hybrid controller while retaining the checkpoint, topology, placements, and every other rule parameter.

D. Observation Design

The observation pipeline is state-based rather than camera-based. Table III shows the complete breakdown. Each agent receives a flat vector containing self-state (position, velocity, heading, alive flag), environment context (time remaining, nearest exit distance and direction, wall proximity, wall shift timer), a fixed-height 360-degree planar horizontal ray fan, last-known-position memory for targets that have left line of sight, and a summary of the two nearest opponents including their relative position, visibility, and threat pressure.

TABLE III
OBSERVATION AND ACTION SPACE PER AGENT ROLE.

Component	Sentinel	Runner
Self state	10	10
Environment context	7	7
Ray perception (6 per ray)	84 (14 rays)	96 (16 rays)
Last-known-position memory	6	6
Opponent summary (2 nearest)	10	10
Total observation size	117	129
Action space	2 continuous	2 continuous

Sentinels cast 14 rays while Runners cast 16 rays. Each ray produces 6 floats: normalized hit distance and one-hot indicators for wall, exit, teammate, opponent, and other object types. This gives Runners wider sensing coverage because their survival depends on detecting threats from all directions.

E. Reward Design

The reward system uses the `v5_active_agents` configuration, which was iteratively developed to address three behavioral failure modes observed during training: (1) free-rider collapse, where one Sentinel learns to capture while others idle, (2) cluster convergence, where all Sentinels chase the same Runner in a single file, and (3) passive evasion, where Runners survive by standing still rather than actively seeking exits.

The reward function combines terminal win/loss signals with individual distance-based shaping for pursuit and evasion [19], anti-clustering penalties, idle penalties, exit-seeking shaping, and a tactical coordination layer that rewards events detected by a dedicated `TrapEventDetector` module. The tactical layer rewards pincer formations, corridor control, exit denial, and dead-end forcing. All weights are loaded from

TABLE IV
REWARD WEIGHTS (V5_ACTIVE_AGENTS CONFIGURATION).

Sentinel Rewards	Weight
Team full capture (win)	+1.0
Timeout or escape (loss)	-1.0
Individual capture bonus	+0.25
Chase progress (per step)	+0.004
Chase regression	-0.003
Cluster penalty	-0.015
Idle penalty	-0.002
Wall loop penalty	-0.002
Exploration visit bonus	+0.001
Orbit stall penalty	-0.003
Runner Rewards	Weight
Team exit success (win)	+1.0
Both captured (loss)	-1.0
Tagged individually	-0.2
Survival step bonus	+0.0001
Exit approach bonus	+0.005
Exit approach regression	-0.002
Evade progress	+0.001
Threat approach penalty	-0.002
Wall loop penalty	-0.001
Exploration visit bonus	+0.001
Orbit stall penalty	-0.003

YAML configuration files at runtime. Table IV lists the final reward weights.

For agent i , the implemented per-step reward can be written as

$$r_{i,t} = r_{i,t}^{\text{terminal}} + \sum_{k \in \mathcal{D}} w_{i,k} g_k(\Delta d_{k,t}) + \sum_{m \in \mathcal{E}} w_{i,m} \phi_m(s_t, \mathbf{a}_t), \quad (8)$$

where $\mathcal{D}$ contains distance-progress terms and $\mathcal{E}$ contains discrete behavioral and tactical events. For a target distance d_t , the shaping gate is

$$g(\Delta d_t) = \mathbb{I}(|\Delta d_t| > \delta) \min\left(1, \frac{|\Delta d_t|}{\kappa}\right), \quad (9)$$

with the sign and coefficient selected by whether the change is desirable. Sentinel chase progress uses $\Delta d_t = d_{t-1} - d_t$. Runner evasion uses $\Delta d_t = d_t - d_{t-1}$ while inside a configured threat radius. Exit approach again uses $d_{t-1} - d_t$. The source defaults use a 0.08 movement dead-zone for pursuit/evasion and 0.05 for exit approach.

The shaping terms are not proven potential-based and may therefore change the optimal policy, not merely accelerate learning. This is a deliberate engineering tradeoff, but it makes reward ablation essential. We evaluate outcomes separately from shaped return so that an increase in accumulated shaping reward cannot be mistaken for improved capture or escape behavior.

F. Logging and Reproducibility

Each training and evaluation run produces four CSV artifacts under a run-specific directory: `episode_log.csv` (per-episode outcomes and totals), `agent_step_log.csv` (per-step positions, actions, and distances), `reward_audit.csv` (per-event reward attribution by agent, category, and magnitude), and

TABLE V
REVIEWER-FACING ARTIFACT MAP IN THE PUBLIC REPOSITORY.

Artifact	Purpose
Unity scripts	Agent control, sensing, reward logic, episode lifecycle, and CSV logging implementation
Configuration files	Trainer, reward, rule, curriculum, ablation, and evaluation manifests
Setup validator	Environment and repository preflight checks before experiments
Policy evaluator	Fixed-checkpoint inference launcher with metadata and log routing
Readiness auditor	Gate-level audit for manuscript, training, evaluation, ablation, and reproducibility status
Unit tests	Evaluation isolation and publication metric calculations
Evaluation summaries	Canonical assist-off 30-cell evaluation summaries and statistics
Diagnostic baselines	Random and geometric-heuristic action-override evaluations with aggregate CSV/JSON outputs
Official summaries	Learning curves, paired ablations, readiness reports, and power analysis outputs

`replay_events.csv` (timestamped capture, escape, wall-shift, and tactical events). This structure ties every reported metric back to raw per-step evidence.

The Python launcher writes a second provenance layer containing the run manifest, seed, exact training or inference command, Git state, configuration snapshots, and SHA-256 hashes. Unity receives the repository root and requested scene through explicit environment variables. Runtime override files select curriculum, rule, and reward configurations. Episode rows repeat the effective scene and configuration names, allowing the audit to compare requested and observed provenance rather than trusting launcher metadata alone.

A run is publication-valid only when all required artifacts are non-empty, the seed and run identifier are known, at least one episode is complete, and the observed scene/rule/reward tuple matches the manifest. This strictness detected an earlier 18-cell matrix in which the standalone silently loaded its default scene and configuration. That matrix is retained with an `INVALIDATED.md` marker but excluded from every reported result. The failure case demonstrates why artifact existence without semantic provenance is insufficient.

The same rule applies to retries. `ML-Agents -force` replaces trainer outputs but does not truncate CSV files written independently by Unity. A corrected training retry was therefore stopped when its V5 rows were found appended to legacy reward rows. The launcher now deletes the complete target run directory before a forced launch and writes an explicit `running` status before starting Unity. The mixed directory is preserved as invalidated audit evidence, never as a source of checkpoints or results.

V. TRAINING AND EVALUATION SETUP

A. Training Algorithm

We train both behaviors with PPO using parameter sharing within each team and asymmetric self-play across teams. All three Sentinels share a single policy network, and both Runners share a separate one. Table VI lists the official maze-stage hyperparameters. The excluded legacy open-arena trainer

TABLE VI
PPO HYPERPARAMETERS (SHARED BY BOTH BEHAVIORS).

Parameter	Static Maze	Dynamic Maze
Learning rate	3.0×10^{-4}	2.5×10^{-4}
Batch size	2048	2048
Buffer size	20480	40960
Entropy (β)	5.0×10^{-3}	4.0×10^{-3}
Clip range (ϵ)	0.2	0.2
GAE λ	0.95	0.95
Discount (γ)	0.99	0.99
Epochs per update	3	4
Hidden units	256	256
Hidden layers	2	2
Time horizon	128	160
Sentinel max steps	1,500,000	1,500,000
Runner max steps	1,500,000	1,500,000

used a 1024 batch and a 64-step horizon and is not part of the four-stage official matrix. The learning rate follows a linear decay schedule over the training run. A prototype MA-POCA YAML exists, but it is not a valid matched baseline: team-group registration, equalized budgets, self-play controls, and independent validation would be required before reporting a MA-POCA comparison. No MA-POCA result or support claim is made here.

For role q , let $r_t(\theta_q) = \pi_{\theta_q}(a_t | o_t) / \pi_{\theta_q^{\text{old}}}(a_t | o_t)$ and let $\hat{A}_{q,t}$ be the generalized advantage estimate. The policy term optimized by PPO is

$$L_q^{\text{clip}}(\theta_q) = \mathbb{E}_t \left[\min \left(r_t(\theta_q) \hat{A}_{q,t}, \text{clip}(r_t(\theta_q), 1 - \epsilon, 1 + \epsilon) \hat{A}_{q,t} \right) \right]. \quad (10)$$

The implementation combines this term with a value loss and entropy bonus. This objective is applied independently to the Sentinel and Runner parameter sets. It does not centralize the critic across opposing roles. Checkpoint transfer initializes both θ_S and θ_R at the next curriculum stage, while optimizer progress and the new stage budget are recorded in run meta-data.

B. Curriculum and Transfer Learning

Training follows a four-stage curriculum with transfer learning, motivated by evidence that staged difficulty helps agents discover coordination skills that are hard to learn from sparse reward alone [18], [29]. Each stage uses `-initialize-from` to bootstrap both role policies from the previous stage’s checkpoint. The network architecture (256 hidden units, 2 layers) remains identical across stages so the parameter tensors are compatible.

Table VII reports the implemented progression from the canonical curriculum YAML. Stage 1 uses a fixed static maze and fixed spawn/exit locations. Stage 2 randomizes spawn and exit locations in the same maze family. Stage 3 introduces one wall shift every 20 seconds. Stage 4 increases the shift frequency to 8 seconds, raises shift intensity from one to three wall changes, and changes the training maze seed from 42 to 84. This is more precise than describing the legacy open-arena plot as the official curriculum.

TABLE VII
IMPLEMENTED FOUR-STAGE CURRICULUM. STEP BUDGETS ARE SENTINEL/RUNNER TRAINER STEPS.

Stage	Topology	Steps	Transfer
1. Static fixed	Static, seed 42	1.5M/1.5M	None
2. Static random	Static, seed 42	1.5M/1.5M	Stage 1
3. Dynamic low	20 s, intensity 1	1.5M/1.5M	Stage 2
4. Dynamic high	8 s, intensity 3	1.5M/1.5M	Stage 3
<i>Evaluation</i>	Seen + 5 held out	None	Stage 4

All four stages use the same `reward_v5_active_agents` definition. A partial 590-episode Stage 1 run was invalidated when an audit found that earlier manifests switched reward definitions between static and dynamic stages. Such a switch would confound curriculum difficulty with reward design. Its first V5 retry was also invalidated because old and new append-only logs shared a run directory. Under simultaneous training, equal limits then froze the faster-accumulating three-agent Sentinel trainer before Runner finished. A 3:2-budget retry avoided that timing error but was also excluded: Sentinel wins fell from 92% in its first 100 episodes to 8% in its last 100 episodes, evidence of latest-policy cycling rather than stable improvement. Opponent-snapshot self-play corrected the simultaneous-update design, but a source-level audit showed that retaining 3:2 limits would instead create a final 500k Sentinel-only tail under the ghost controller’s equal-step turns. That partial run was also invalidated. The final configuration combines snapshot self-play with equal per-policy sample limits. The launcher refuses official training unless reward, clean-retry, role-budget, and self-play audits pass.

ML-Agents counts one wrapped-trainer step for each acting-agent transition. Its ghost controller changes the learning team after an equal number of wrapped-trainer steps and is designed to give every team the same trainer-step count. We therefore use 1.5M transition samples for each role per stage, so both optimizers receive equal sample budgets and finish self-play together. With role-level parameter sharing, these samples correspond to 0.5M Sentinel team decision cycles and 0.75M Runner team decision cycles because three Sentinels act per cycle but only two Runners do. We disclose both units rather than calling them equivalent. The matched direct-dynamic control uses 6M samples per role, exactly the aggregate four-stage curriculum budget.

C. Evaluation Protocol

We evaluate trained policies in inference mode using fixed ONNX checkpoints. The strengthened publication protocol uses five policy seeds (42, 101, 202, 606, 707), one seen dynamic-maze split, and at least five held-out unseen layout seeds (101, 202, 303, 404, 505). Each policy-layout cell targets 100 completed episodes. A 30-minute wall-clock cap is a failure timeout rather than the sampling unit. All evaluation runs use deterministic action selection with seeded spawn/episode randomization and must produce run-scoped

raw logs: episode_log.csv, agent_step_log.csv, reward_audit.csv, and replay_events.csv. A run is excluded from publication tables if the episode target, any required log, metadata file, or KPI summary is missing.

The seen and held-out rule files are matched on agent speed, capture radius, timeout, wall-shift interval and intensity, sensing, zero action assistance, and all enabled randomization switches. Only the split identifier, topology seed, and placement/event seed differ. A machine-readable control audit compares the nested YAML fields before launch and rejects a matrix with any unregistered difference. The action-assist-on and dynamic-wall-off matrices use a parallel single-intervention audit: the former may change only the two assist coefficients, and the latter may change only dynamic_walls.enabled.

The evaluation protocol reports outcome metrics (Sentinel win rate, Runner win rate, escape rate, survival time, and capture times), coordination metrics (pincer episode rate, trap frequency, corridor-block events, exit-denial events, Sentinel spread, and Runner separation), path metrics (path efficiency and route-change proxies after wall shifts), and ablation deltas. Primary results use zero action assistance. The paired assist-on matrix quantifies how much the optional hybrid controller changes those results.

D. Metric Definitions

For event type m , let $C_{k,m}$ be its count in episode k . We distinguish an episode prevalence from an event intensity:

$$\hat{P}_m = \frac{1}{n} \sum_{k=1}^n \mathbb{I}[C_{k,m} > 0], \quad \bar{C}_m = \frac{1}{n} \sum_{k=1}^n C_{k,m}. \quad (11)$$

Calling both quantities a “rate” would be ambiguous, so exported summaries use names such as pincer_episode_rate and pincer_events_per_episode. Trap success is conditional:

$$\hat{P}_{\text{trap} \rightarrow \text{win}} = \frac{\sum_k \mathbb{I}[C_{k,\text{trap}} > 0 \wedge y_k = S]}{\sum_k \mathbb{I}[C_{k,\text{trap}} > 0]}. \quad (12)$$

At step t , Sentinel spread is the mean pairwise distance among active Sentinels,

$$D_S(t) = \left(\frac{|\mathcal{I}_S^t|}{2} \right)^{-1} \sum_{i < j} \|x_i(t) - x_j(t)\|_2, \quad (13)$$

with an analogous Runner separation $D_R(t)$. Each spread is averaged over snapshots within an episode and then equally across episodes. Path length is $L_i = \sum_t \|x_i(t) - x_i(t-1)\|_2$, with positions keyed by episode and agent so resets are not counted as motion. Capture times come from replay-event time_seconds. Runner survival ends at that Runner’s capture or at episode termination and is averaged over runner-episodes. For a wall shift at τ and one-second window w , route response is

$$\begin{aligned} \theta_{j,\tau} &= \cos^{-1} \left(\frac{u^- \cdot u^+}{\|u^-\| \|u^+\|} \right), \\ u^- &= x_j(\tau) - x_j(\tau - w), \\ u^+ &= x_j(\tau + w) - x_j(\tau). \end{aligned} \quad (14)$$

We report mean absolute deflection and the fraction with $\theta_{j,\tau} \geq 45^\circ$. Windows with less than 0.05 m displacement are excluded as directionally undefined.

E. Statistical Analysis Plan

The independent training seed is the unit of replication for policy-level claims. Episodes from the same checkpoint are not treated as independent policy replicates. We report each seed separately, then the mean and standard deviation across trained policies. Binary outcomes receive Wilson intervals within a policy/layout cell. For aggregate comparisons, a hierarchical bootstrap resamples policy seeds, then layouts, then episodes to preserve the nested design. Training uses the fixed max_steps budgets in Table VII. The YAML episode-count and win-rate thresholds are audited post hoc and do not terminate ML-Agents training.

Training curves are exported directly from each audited ML-Agents TensorBoard event file. Curves include cumulative reward, episode length, self-play ELO, entropy, policy loss, and value loss. They are aligned by role, curriculum stage, and normalized within-stage trainer progress. No run is admitted unless its completion audit passes. At every matched scalar step, the plot reports per-seed traces, the arithmetic mean, and a two-sided 95% Student- t interval across independent policy seeds. With five seeds the interval remains descriptive and potentially wide. Smoothing is limited to the summary aggregation already performed by ML-Agents.

The full condition contains $5 \times 6 \times 100 = 3,000$ completed evaluation episodes. Applying the same budget to five ablations yields a planned minimum of 18,000 completed episodes after all retrained checkpoints are available. This count does not create 18,000 independent policy replicates: inference about training variability remains limited to five policy seeds. The larger within-cell count instead stabilizes per-topology outcome and timing estimates.

For metric observation $z_{s,l,k}$ from policy seed s , layout l , and episode k , the layout-balanced estimator is

$$\hat{M} = \frac{1}{|S|} \sum_{s \in S} \frac{1}{|\mathcal{L}|} \sum_{l \in \mathcal{L}} \left(\frac{1}{n_{s,l}} \sum_{k=1}^{n_{s,l}} z_{s,l,k} \right). \quad (15)$$

This prevents fast-terminating layouts from receiving more weight merely because they produce more episodes. The fixed-count protocol makes episode-weighted and layout-balanced estimates closer, but both are exported. Ablation effects are paired by training seed and evaluation layout:

$$\Delta_{a,s,l} = M_{\text{full},s,l} - M_{a,s,l}, \quad (16)$$

where a denotes an ablated component, s a policy seed, and l a layout. We report the mean paired effect, a 95% bootstrap interval, and the raw per-seed values. With only five policy seeds, interval width and effect consistency matter more than a binary significance threshold. For descriptive scale we also report the paired standardized effect

$$d_z = \bar{\Delta}_a / s(\Delta_a), \quad (17)$$

along with the fraction of matched seed-topology cells having the modal effect sign. We retain the raw metric units because

d_z is unstable with a small policy-seed sample or when paired variance is near zero. The corresponding two-sided paired- t power function is

$$\pi(d, n) = \Pr(|T_{n-1, d\sqrt{n}}| > t_{1-\alpha/2, n-1}), \quad (18)$$

where $T_{\nu, \lambda}$ is noncentral t . At $\alpha = 0.05$, $n = 5$, and 80% power, numerical inversion gives $|d_z| \approx 1.68$. Thus even five seeds detect only large effects: more episodes improve within-cell precision but not policy-level power, and a nonsignificant seed test is not evidence of no effect. The exported power audit also reports thresholds of 3.26 for three seeds and 1.00 for ten.

For outcome metric M , the topology generalization gap for policy seed s is

$$G_s = M_{s, \text{seen}} - \frac{1}{5} \sum_{l \in \mathcal{L}_{\text{held}}} M_{s, l}. \quad (19)$$

For Sentinel win rate, $G_s > 0$ means Sentinel success decreases on held-out topologies. For Runner escape rate, $G_s > 0$ means Runner escape success decreases. We state the role perspective for every gap rather than assigning one global “favorable” direction. We report the five held-out values alongside G_s so that a favorable mean cannot hide a topology-specific failure. The statistics pipeline bootstraps each G_s by resampling the seen episodes and then held-out layouts and their episodes. It also exports a balanced two-way decomposition of policy-seed variation, layout variation, interaction, and finite-cell error from the cell means. No test treats the five layouts as five independently trained policies.

For held-out cell mean $\bar{z}_{s, l}$, the descriptive random-effects decomposition is

$$\bar{z}_{s, l} = \mu + a_s + b_l + u_{s, l}, \quad \text{Var}(\bar{z}) = \sigma_a^2 + \sigma_b^2 + \sigma_u^2, \quad (20)$$

where a_s is policy-seed variation, b_l is layout variation, and $u_{s, l}$ combines seed and layout interaction with finite-cell estimation error. Method-of-moments components are truncated at zero and reported as variance fractions. With only five policy seeds they are descriptive rather than precise population estimates.

Memory, tactical-reward, and curriculum ablations alter the learning problem and therefore require paired retraining for every policy seed. Dynamic-wall and action-assist ablations are deployment-time interventions on the same checkpoints. Pairing the policy and episode/topology seeds isolates the immediate contribution of those mechanisms. We do not mix these two ablation types in one effect estimate.

The machine-readable suite pre-registers signs under the convention *canonical full minus condition*: shorter reacquisition with memory, higher tactical-event prevalence with tactical shaping, a smaller absolute transfer gap after curriculum training, higher Sentinel win rate with assist on, and stronger route deflection with dynamic walls. Metrics outside these hypotheses are explicitly exploratory. Disagreement with a registered sign is reported as a negative result rather than relabeled post hoc.

The memory intervention is a system-level ablation rather than a neural-input-only mask. It writes zeros into the fixed




Fig. 4. Audited five-seed official curriculum reward curves exported from TensorBoard scalars. The official gate uses fixed trainer step budgets. The plotted curves are diagnostic summaries of completed training.




Fig. 5. Audited five-seed self-play ELO traces. Stable or increasing ELO does not imply Nash convergence. It is used as a training diagnostic.

six-dimensional memory observation, stops last-known-target updates, and prevents the pursuit-assist controller from falling back to a remembered target when no opponent is visible. This preserves observation dimensionality and network compatibility, and any measured effect is attributed to the complete explicit-memory subsystem.

VI. EMPIRICAL RESULTS

Evidence scope. This section reports the completed publication-evaluation evidence used for the paper: five independently trained curriculum policy seeds, fixed-count seen and held-out topology evaluation, low-cost control policies, and two paired deployment interventions. Retraining interventions for memory, tactical reward shaping, and direct dynamic training are specified in the ablation protocol but are treated as future extensions rather than evidence for the claims below.

The official five-seed curriculum baseline reached 20/20 audited training runs. Figure 4 summarizes the exported TensorBoard reward traces for the five independent policy seeds. The curve export contains 28,320 scalar rows and is generated from `results/official_summary/training_curves/`. It is an audited learning-curve artifact, not a stopping criterion: all official stages use fixed trainer step budgets.

Figure 5 reports the corresponding self-play ELO traces. ELO is useful for detecting cycling and relative opponent difficulty, but it is not treated as a game-theoretic equilibrium certificate.

A. Canonical Assist-Off Evaluation

The canonical evaluation uses the five trained curriculum policy seeds (42, 101, 202, 606, 707), one seen topology, five held-out topologies, and a target of 100 completed episodes per

TABLE VIII
PRE-REGISTERED ABLATION QUESTIONS AND ACCEPTANCE EVIDENCE. EACH CONDITION USES POLICY SEEDS 42, 101, 202, 606, AND 707 AND THE SAME HELD-OUT LAYOUTS.

Ablation	Controlled change	Primary metrics	Interpretation
Memory off	Retrain with the last-known-position subsystem disabled	Win/escape rate, capture time, reacquisition delay	Tests whether explicit system memory helps under occlusion
Tactical reward off	Retrain with trap/pincer/corridor/exit-denial rewards zeroed	Outcome rates and event prevalence/intensity	Separates coordination behavior from reward accounting
Dynamic walls off	Evaluate the same policies with matched topology/placement seeds	Survival, path length, route-change proxy	Measures the cost of within-episode topology changes
Direct dynamic training	Retrain without checkpoint transfer and match total environment steps	Learning curve, final outcome rates, seed variance	Tests the curriculum hypothesis
Action assist on	Evaluate the same policies with Sentinel pursuit blend restored to 0.15	All outcome and coordination metrics	Quantifies non-policy control contribution relative to the canonical assist-off evaluation

TABLE IX
CANONICAL ASSIST-OFF EVALUATION OVER FIVE POLICY SEEDS. VALUES ARE MEANS OVER POLICY-TOPOLOGY CELLS. THE SECOND NUMBER IS THE ACROSS-CELL STANDARD DEVIATION.

Metric	Seen	Five held-out topologies
Cells / episodes	5 / 501	25 / 2,501
Sentinel win rate	41.5% $\pm$ 5.6	42.2% $\pm$ 8.2
Runner win rate	58.5% $\pm$ 5.6	57.8% $\pm$ 8.2
Escape rate	52.9% $\pm$ 8.3	53.5% $\pm$ 9.6
Full capture time	23.70 $\pm$ 5.09 s	22.79 $\pm$ 4.04 s
Pincer episode rate	22.6% $\pm$ 4.5	20.3% $\pm$ 4.5
Corridor block rate	4.0% $\pm$ 1.8	2.5% $\pm$ 1.8
Exit denial rate	42.7% $\pm$ 3.3	39.2% $\pm$ 4.7
Trap success rate	49.4% $\pm$ 4.7	50.4% $\pm$ 8.2

policy-topology cell. This yields 30 completed cells and 3,002 episodes because two cells slightly exceeded the target before shutdown. Every reported row comes from the canonical evaluation aggregate and passed the strict publication-evaluation audit.

Table IX changes the interpretation of the earlier legacy story. Under canonical assist-off evaluation, the task is not Sentinel-dominated. Runner wins are more common than Sentinel wins on both the seen and held-out splits, and escapes account for roughly half of all episodes. This is evidence of a playable asymmetric benchmark with meaningful evasion, not evidence that the pursuer policy has solved the task. The held-out split is not materially worse than the seen split at this aggregate level, but the policy-seed variation is substantial: mean Sentinel win rate across the six layouts ranges from 31.6% for seed 101 to 50.8% for seed 606. This is why the paper reports per-seed results and uses the bounded term held-out topology transfer rather than broad zero-shot generalization.

B. Lightweight Evaluation-Only Baselines

To answer the reviewer request for a state-of-practice comparison without claiming an unearned retrained SOTA baseline, we added two evaluation-only controllers to the Unity runtime. The baseline selector is controlled by `LABYRINTH_BASELINE_POLICY` or by the runtime over-

ride file written by `scripts/evaluate_policy.py`. In random mode, both teams receive deterministic pseudo-random continuous actions keyed by episode, step, agent identifier, and seed. In heuristic mode, Sentinels greedily steer toward the nearest active Runner, while Runners blend steering away from the nearest Sentinel with steering toward the nearest exit. Both use local wall avoidance. The ML-Agents checkpoint path is still supplied only to maintain the standard inference handshake, but the Unity action override ignores the neural policy outputs.

Table X reports a 12-cell low-cost comparison using one evaluation seed and 25 target episodes per split. These rows provide control-policy context rather than a matched-compute learned-baseline leaderboard. They show that the learned PPO policy is substantially stronger than random pursuit on held-out capture time and coordination prevalence, but a hand-coded geometric controller still captures faster and wins more often. The defensible conclusion is therefore specific: Labyrinth Breach is a challenging benchmark in which the current PPO policy is competent but not dominant.

C. Paired Deployment Interventions

Two paired deployment-time interventions are reported: action-assist on/off and dynamic-wall on/off. The convention is canonical full condition minus ablated condition. Thus, for *action assist on*, a negative Sentinel win delta means the canonical assist-off policy won less often than the assist-on hybrid controller. For *dynamic wall off*, a positive tactical event delta means the dynamic-wall condition produced more of that event than the static-wall intervention.

The completed action-assist intervention does not show a reliable win-rate benefit under the paired held-out analysis: the bootstrap interval spans zero. It does, however, change some spatial behavior. Canonical assist-off runs show lower Sentinel spread under the full-minus-assist-on convention. This supports disclosing assist settings but does not justify claiming that geometric assistance is necessary for success.

The dynamic-wall intervention has clearer tactical effects. With dynamic walls enabled, pincer episodes, exit denial, and trap episodes are higher on held-out topologies, and stall-step fraction is lower. The Sentinel win-rate delta still spans

TABLE X

LOW-COST EVALUATION-ONLY BASELINES. LEARNED PPO USES THE CANONICAL FIVE-POLICY-SEED ASSIST-OFF EVALUATION FROM TABLE IX. RANDOM AND HEURISTIC ROWS USE ONE EVALUATION SEED, SIX SPLITS, AND 25 TARGET EPISODES PER SPLIT. HELD-OUT VALUES ARE MEAN $\pm$ STANDARD DEVIATION OVER FIVE LAYOUT CELLS.

Controller	Evidence	S win seen	S win held out	Escape seen	Escape held out	Full cap. held out	Pincer held out
Learned PPO	5 policy seeds, 30 cells, 3,002 episodes	41.5 $\pm$ 5.6	42.2 $\pm$ 8.2	52.9 $\pm$ 8.3	53.5 $\pm$ 9.6	22.79 $\pm$ 4.04 s	20.3 $\pm$ 4.5
Random actions	1 eval seed, 6 cells, 150 episodes	32.0	27.2 $\pm$ 4.4	48.0	44.8 $\pm$ 4.4	103.03 $\pm$ 13.23 s	8.8 $\pm$ 4.4
Geometric heuristic	1 eval seed, 6 cells, 152 episodes	73.1	68.3 $\pm$ 9.7	26.9	30.1 $\pm$ 9.5	8.74 $\pm$ 0.59 s	26.3 $\pm$ 13.0

TABLE XI

PAIRED DEPLOYMENT INTERVENTIONS ON HELD-OUT TOPOLOGIES. DELTAS USE THE FULL-MINUS-CONDITION CONVENTION AND ARE PAIRED BY POLICY SEED AND TOPOLOGY.

Ablation / metric	Delta	d_z	Boot. low	Boot. high
Assist on: Sentinel win	-1.26 pp	-0.17	-6.88 pp	4.44 pp
Assist on: escape	1.82 pp	0.25	-3.64 pp	7.38 pp
Assist on: Sentinel spread	-0.153 m	-0.48	-0.290	-0.011
Assist on: pincer rate	0.91 pp	0.16	-2.09 pp	4.50 pp
Wall off: Sentinel win	2.06 pp	0.29	-2.74 pp	7.22 pp
Wall off: pincer rate	2.75 pp	0.63	1.08 pp	4.52 pp
Wall off: exit denial	2.34 pp	0.45	0.20 pp	4.58 pp
Wall off: stall fraction	-6.46 pp	-1.21	-8.77 pp	-3.77 pp
Wall off: trap episode rate	3.16 pp	0.76	1.08 pp	5.40 pp

zero, so the defensible claim is behavioral: dynamic topology changes the coordination-event profile more clearly than it changes terminal win rate in the completed sample.

D. Interpretability from Existing Artifacts

The present manuscript does not report SHAP values or gradient saliency maps. Instead, it provides three layers of model interpretation that can be generated from the existing evidence pack. First, `reward_audit.csv` identifies which terminal, shaping, penalty, and tactical events contributed reward at each episode stage. This makes it possible to separate a real outcome change from a shaped-return increase. Second, `replay_events.csv` exposes behavioral mechanisms such as pincer pressure, corridor blocks, exit denial, captures, escapes, and wall shifts. Third, the paired action-assist and dynamic-wall interventions act as system-level counterfactual probes: the same policy, topology seeds, and placement seeds are replayed while changing one mechanism.

Under this interpretation lens, the current evidence is mixed and useful. The assist-on probe does not show a reliable Sentinel win-rate gain, which weakens any claim that the learned policy merely depends on geometric pursuit assistance. The wall-off probe changes tactical-event prevalence more clearly than terminal outcomes, suggesting that dynamic topology affects how agents coordinate even when it does not yet produce a statistically clear win-rate shift. This style of explanation is coarser than feature attribution, but it is well matched to MARL: it interprets behavior at the event, team, and mechanism levels where reviewers can connect policy outputs to pursuit-evasion concepts. Future SHAP-style analysis should group inputs by semantic channels, including self state, wall rays, exit rays, opponent rays, teammate rays,

memory, and time context, rather than treating each scalar feature independently.

E. Registered Extensions

Three additional ablation families are registered for extended study: memory-off, tactical-reward-off, and direct-dynamic training. These are retraining interventions and cannot be inferred from the canonical checkpoints. They are therefore excluded from the primary evidence claims in this article. This separation prevents deployment-time interventions and retraining interventions from being mixed in one effect estimate.

VII. DEVELOPMENT PROVENANCE AND EXCLUDED DIAGNOSTICS

The official empirical claims in this article use only the v5 five-seed training matrix, the 30-cell canonical assist-off evaluation, the lightweight control policies, and the two paired deployment interventions reported above. Earlier development runs are archived for transparency because they shaped the final evaluation design. They are not counted as official evidence.

TABLE XII

ARCHIVED DEVELOPMENT DIAGNOSTICS AND THEIR EFFECT ON THE FINAL PROTOCOL.

Archived item	Diagnostic finding	Protocol change
Three-phase training trace v4	Sentinel dominance decreased after dynamic walls were introduced	Official results use the four-stage v5 curriculum and five independent policy seeds
Fixed-checkpoint held-out diagnostic	Seeded configurations changed placements more than topology	Official evaluation separates policy seed, topology seed, and episode seed
Assist-on checkpoint diagnostic	Geometric pursuit assistance could affect Sentinel behavior	Canonical results are assist-off and assist-on is reported as a paired intervention
Reward decomposition audit	Terminal reward remained visible beside shaping signals	Reward events are logged through <code>reward_audit.csv</code> and tactical reward is registered for retraining study

The earlier fixed-checkpoint diagnostic used a checkpoint named `v5_dynamicmaze`, but episode provenance showed that inference loaded the legacy `reward_dynamicmaze_memory_v4.yaml` file, retained

Sentinel action assistance at $\rho_S = 0.15$, reused one hard-coded unseen topology, and combined topology changes with movement and wall-rule changes. The final protocol resolves these confounds by using independently trained policy seeds, explicit topology seeds, fixed completed-episode targets, matched non-topology controls, and canonical assist-off inference. This audit trail is included to make the evolution from exploratory runs to controlled evidence transparent.

A. Archived Diagnostic Figures

The legacy figures below are retained only as development provenance. They show why the final study emphasizes curriculum transfer, reward provenance, assist disclosure, and controlled held-out topology evaluation. They should be read as engineering diagnostics rather than as official performance claims.

VIII. DISCUSSION

The completed evidence supports three bounded findings.

Curriculum training produced stable deployable policies. The four-stage curriculum with transfer learning produced five audited policy seeds that complete the canonical evaluation matrix and maintain comparable aggregate performance on held-out layouts. The direct-dynamic extension remains useful for isolating the curriculum effect from total training exposure.

Custom reward shaping targets specific failure modes.

The iterative development of the `v5_active_agents` reward configuration was driven by three observed failure modes:

- *Free-rider collapse:* In early training, one Sentinel learned to capture while others remained idle to collect the shared team reward. The `idle_penalty` ($-0.002/\text{step}$) and increased `individual_capture_bonus` ($+0.25$) were added to reduce this behavior.
- *Cluster convergence:* All three Sentinels would chase the same Runner in single file, making evasion trivial. The `cluster_penalty` (-0.015) encourages spatial separation. The official results quantify this through Sentinel spread and pincer/corridor metrics.
- *Passive evasion:* Runners learned to survive by standing still, avoiding threat-approach penalties but never seeking exits. The `exit_approach_bonus` ($+0.005$) and `exit_approach_regression` (-0.002) were added to encourage goal-directed escape behaviour.

Held-out topology transfer is measurable but not overstated. The official protocol tests the upgraded procedural generator using independent policy seeds and five topology seeds, with topology identity recorded in every episode row. The resulting aggregate held-out performance is comparable to the seen split, but policy-seed variation remains large enough that the paper avoids broad zero-shot claims.

Compared to MatrixWorld [9], which provides a grid-based pursuit-evasion benchmark, Labyrinth Breach operates in continuous 3D space with asymmetric teams, dynamic walls, and exit-based objectives. Compared to ViPER [13],




Fig. 6. Archived three-phase cumulative reward used during development.




Fig. 7. Archived reward decomposition motivating reward-audit logging.

which focuses on visibility-based pursuit, our environment adds team coordination pressure through the 3v2 asymmetry and tactical reward terms.

Contribution scope. The article is framed as a reproducible simulation benchmark and empirical study for multidisciplinary robotics and AI readers. The novelty is the combination of asymmetric learned teams, dynamic topology, exit-directed evasion, partial observation, explicit reward provenance, and artifact-gated evaluation. Mechanism-level claims are limited to the evidence actually reported: canonical evaluation, control-policy diagnostics, action-assist intervention, and dynamic-wall intervention.

IX. CODE, DATA, AND REPRODUCIBILITY

The source repository contains the Unity project, ML-Agents trainer manifests, reward/rule/curriculum YAML files, Python launchers, audit scripts, unit tests, manuscript source, and generated summary artifacts. The public URL is given in the abstract footnote. The intended reproduction path is documented in `docs/reproducibility_guide.md`: build the Unity 6000.0.40f1 standalone, run the prelaunch validation gates, resume the official curriculum/evaluation/ablation launchers, export KPI summaries, re-run readiness audits, and compile the paper. The Python environment is specified by `environment.yml` and `requirements.txt`. The core versions are Unity ML-Agents package `com.unity.ml-agents@4.0.0`, Python 3.10.12, `mlagents==1.1.0`, and `torch==2.2.1`.

The current PDF is generated from the journal L^AT_EX source under `paper/`. Official evaluation results are read from the publication-evaluation aggregate directory. Learning curves and paired deployment interventions are read from the official-summary directory. The codebase includes unit tests under `tests/`. The current-evidence snapshot script exports a concise JSON, CSV, and Markdown summary for audit review. The evidence-pack builder creates a hashed archive with `artifact_manifest.json` and SHA256SUMS so every included file can be checked independently.

TABLE XIII
CLAIM-TO-ARTIFACT TRACEABILITY FOR THE CURRENT MANUSCRIPT.

Manuscript item	Source artifact	Regeneration command
Five-seed learning curves	results/official_summary/training_curves/	scripts/analyze_official_training_curves.py
Canonical evaluation table	results/publication_eval_official/aggregate/	scripts/run_publication_eval_matrix.py
Random and heuristic diagnostics	results/lightweight_baselines/aggregate/	scripts/run_lightweight_baseline_eval.py
Action-assist intervention	results/official_summary/ablations/action_assist_on/	scripts/analyze_paired_ablation.py
Dynamic-wall intervention	results/official_summary/ablations/dynamic_wall_off/	scripts/analyze_paired_ablation.py
Readiness status	results/official_summary/publication_readiness.json	scripts/audit_publication_readiness.py
Current evidence snapshot	results/official_summary/current_evidence_snapshot.*	scripts/build_current_evidence_snapshot.py




Fig. 8. Archived win-rate trace used to motivate fixed-count evaluation.




Fig. 9. Archived penalty traces motivating coordination and safety metrics.

This revision also includes a random-action and geometric-heuristic baseline harness for evaluation-only controls. Those rows are generated through the same fixed-checkpoint ML-Agents launch path and KPI summarizer, with Unity overriding policy actions at runtime. Their retained evidence level is KPI/metadata summaries plus aggregate CSV/JSON outputs, so they are reported as control-policy context rather than a formal matched-compute SOTA comparison. A learned MARL baseline under equalized compute is a natural extension.

X. LIMITATIONS AND THREATS TO VALIDITY

Several limitations remain. First, the study is simulation-only and does not include physical robot validation. Second, wall collision correctness under all dynamic shift configurations has not been exhaustively validated. Rare edge cases during wall transitions could allow agents to clip through geometry. Third, the tactical reward shaping terms add complexity to the reward landscape and could be exploited in ways not yet detected. Fourth, the current architecture uses last-known-position memory rather than recurrent neural memory,




Fig. 10. Archived shaping traces motivating the tactical-reward extension.




Fig. 11. Archived per-phase summary used to design the official curriculum matrix.

which may limit performance in highly dynamic environments. Fifth, the controller includes optional action-assist blending. Canonical results are assist-off, and the assist-on intervention is reported separately.

Future work should prioritize retraining interventions for memory, tactical reward shaping, and direct dynamic training, a matched-compute learned MARL baseline, recurrent or belief-state memory, and stress tests for wall-transition safety. SHAP-style attribution is also useful, but it should group features by semantic channels before producing visual explanations.

XI. CONCLUSION

We presented Labyrinth Breach, a reproducible Unity ML-Agents benchmark for asymmetric multi-agent pursuit-evasion with dynamic topology, partial observation, exit-directed evasion, and tactical reward shaping. The evidence includes audited five-seed training curves, a 30-cell canonical assist-off evaluation, control-policy baselines, and two paired deployment interventions. The benchmark remains challenging:

TABLE XIV
MAIN VALIDITY THREATS AND CURRENT MITIGATION.

Threat	Mitigation in this version
Simulation scope	Claims are limited to Unity simulation and no hardware transfer is claimed
Seed power	Five policy seeds are reported, but the power audit still marks small effects as hard to resolve
Baseline scope	Random and heuristic controllers are diagnostic controls rather than learned SOTA baselines
Reward confounding	Reward events are logged and tactical-reward retraining is registered as an extension
Assisted control	Canonical results are assist-off and the assist-on intervention is reported separately
Topology transfer	Held-out topology seeds are recorded and per-topology evaluation is preserved

Runner wins and escapes are common, and a geometric heuristic captures faster than the current PPO policy. These results support Labyrinth Breach as a controlled benchmark for studying asymmetric MARL behavior under changing topology and artifact-gated evaluation.

REFERENCES

- [1] AI641: AI for Robotics, “Projects,” course handout, Project 2: Reinforcement Learning-based multi-agent pursuit-evasion.
- [2] R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch, “Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments,” in *Advances in Neural Information Processing Systems*, 2017.
- [3] T. Rashid, M. Samvelyan, C. S. de Witt, G. Farquhar, J. Foerster, and S. Whiteson, “QMIX: Monotonic Value Function Factorisation for Deep Multi-Agent Reinforcement Learning,” *Journal of Machine Learning Research*, vol. 21, no. 178, pp. 1–51, 2020.
- [4] C. Yu et al., “The Surprising Effectiveness of PPO in Cooperative Multi-Agent Games,” in *Advances in Neural Information Processing Systems*, 2022.
- [5] T. Verma, F. M. Zegers, and P. P. Khargonekar, “Multi-Agent Pursuer-Evader Learning using Situation Report,” arXiv preprint arXiv:1910.07780, 2019.
- [6] M. Zhou, Z. Liu, P. Sui, Y. Li, and Y. Y. Chung, “Learning Implicit Credit Assignment for Cooperative Multi-Agent Reinforcement Learning,” in *Advances in Neural Information Processing Systems*, 2020.
- [7] W. Chen, W. Li, X. Liu, S. Yang, and Y. Gao, “Learning Explicit Credit Assignment for Cooperative Multi-Agent Reinforcement Learning via Polarization Policy Gradient,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, no. 10, pp. 11542–11550, 2023, doi: 10.1609/aaai.v37i10.26364.
- [8] O. Gonultas and V. Isler, “Learning to Play Pursuit-Evasion with Dynamic and Sensor Constraints,” arXiv preprint arXiv:2405.05372, 2024.
- [9] L. Sun, Y.-C. Chang, C. Lyu, C.-T. Lin, and Y. Shi, “MatrixWorld: A Pursuit-Evasion Platform for Safe Multi-Agent Coordination and Autocurricula,” arXiv preprint arXiv:2307.14854, 2023.
- [10] M. Kouzeghar, Y. Song, M. Meghjani, and R. Bouffanais, “Multi-Target Pursuit by a Decentralized Heterogeneous UAV Swarm using Deep Multi-Agent Reinforcement Learning,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2023.
- [11] J. Chen, G. Li, C. Yu, X. Yang, B. Xu, H. Yang, and Y. Wang, “A Dual Curriculum Learning Framework for Multi-UAV Pursuit-Evasion in Diverse Environments,” arXiv preprint arXiv:2312.12255, 2023.
- [12] Y. Zheng, Y. Zhang, C. Zhao, H. Yang, T. Li, Q. Ouyang, and Y. Chen, “Faster Target Encirclement with Utilization of Obstacles via Multi-Agent Reinforcement Learning,” in *Proc. Asian Conf. Mach. Learn., PMLR*, vol. 222, pp. 1731–1746, 2024.
- [13] Y. Wang, Y. Cao, J. Chiun, S. Koley, M. Pham, and G. A. Sartoretti, “ViPER: Visibility-based Pursuit-Evasion via Reinforcement Learning,” in *Proceedings of the 8th Conference on Robot Learning*, PMLR, vol. 270, pp. 4188–4200, 2025.
- [14] R. Lu, P. Zhang, R. Shi, Y. Zhu, D. Zhao, Y. Liu, D. Wang, and C. Alippi, “Equilibrium Policy Generalization: A Reinforcement Learning Framework for Cross-Graph Zero-Shot Generalization in Pursuit-Evasion Games,” in *Advances in Neural Information Processing Systems*, 2025.
- [15] R. Lu, R. Shi, Y. Zhu, and D. Zhao, “R2PS: Worst-Case Robust Real-Time Pursuit Strategies under Partial Observability,” in *Proc. Int. Conf. Learn. Representations (ICLR)*, 2026.
- [16] F. Yan, J. Jiang, K. Di, Y. Jiang, and Z. Hao, “Multiagent Pursuit-Evasion Problem with the Pursuers Moving at Uncertain Speeds,” *Journal of Intelligent & Robotic Systems*, vol. 95, no. 1, pp. 119–135, 2019.
- [17] J. Foerster, G. Farquhar, T. Afouras, V. Nardelli, and S. Whiteson, “Counterfactual Multi-Agent Policy Gradients,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, no. 1, pp. 2974–2982, 2018, doi: 10.1609/aaai.v32i1.11794.
- [18] R. Bhati, S. K. Gottipati, C. Mars, and M. E. Taylor, “Curriculum Learning for Cooperation in Multi-Agent Reinforcement Learning,” in *Second Agent Learning in Open-Endedness Workshop at NeurIPS*, 2023.
- [19] H. U. Sheikh and L. Bölöni, “Multi-Agent Reinforcement Learning for Problems with Combined Individual and Team Reward,” arXiv preprint arXiv:2003.10598, 2020.
- [20] Y. Zhang, M. Ding, J. Zhang, Q. Yang, G. Shi, M. Lu, and F. Jiang, “Multi-UAV Pursuit-Evasion Gaming Based on PSO-M3DDPG Schemes,” *Complex & Intelligent Systems*, vol. 10, pp. 6867–6883, 2024, doi: 10.1007/s40747-024-01504-1.
- [21] W. Du, T. Guo, J. Chen, B. Li, G. Zhu, and X. Cao, “Cooperative Pursuit of Unauthorized UAVs in Urban Airspace via Multi-Agent Reinforcement Learning,” *Transportation Research Part C: Emerging Technologies*, vol. 128, Art. no. 103122, 2021, doi: 10.1016/j.trc.2021.103122.
- [22] Z. Peng, G. Wu, B. Luo, and L. Wang, “Multi-UAV Cooperative Pursuit Strategy with Limited Visual Field in Urban Airspace: A Multi-Agent Reinforcement Learning Approach,” *IEEE/CAA Journal of Automatica Sinica*, vol. 12, no. 7, pp. 1350–1367, 2025, doi: 10.1109/JAS.2024.124965.
- [23] X. Qu, C. Li, Y. Jiang, F. Long, and R. Zhang, “Cooperative Pursuit of Unmanned Surface Vehicles Using Multi-Agent Reinforcement Learning,” *Journal of Shanghai Jiaotong University (Science)*, vol. 31, pp. 187–194, 2026, doi: 10.1007/s12204-025-2816-6.
- [24] S. Xu and Z. Dang, “Emergent Behaviors in Multiagent Pursuit Evasion Games within a Bounded 2D Grid World,” *Scientific Reports*, vol. 15, Art. no. 29376, 2025, doi: 10.1038/s41598-025-15057-x.
- [25] E. Cetin, C. Barrado, E. Salami, and E. Pastor, “Analyzing Deep Reinforcement Learning Model Decisions with Shapley Additive Explanations for Counter Drone Operations,” *Applied Intelligence*, vol. 54, pp. 12095–12111, 2024, doi: 10.1007/s10489-024-05733-2.
- [26] L. Sheng, X. Chen, Z. Pan, F. Cai, and H. Chen, “From Individual Decisions to Team Emergence: A Survey on Explainable Cooperative Multi-Agent Reinforcement Learning,” *Artificial Intelligence Review*, 2026, doi: 10.1007/s10462-026-11598-3.
- [27] P. K. Sharma, E. Zaroukian, D. E. Asher, and B. Howell, “Emergent Behaviors in Multi-Agent Target Acquisition,” arXiv preprint arXiv:2212.07891, 2022.

- [28] C. de Souza Jr., R. Newbury, A. Cosgun, P. Castillo, B. Vidolov, and D. Kulić, “Decentralized Multi-Agent Pursuit using Deep Reinforcement Learning,” *IEEE Robotics and Automation Letters*, vol. 6, no. 3, pp. 4552–4559, 2021.
- [29] N. A. Grupen, D. D. Lee, and B. Selman, “Multi-Agent Curricula and Emergent Implicit Signaling,” in *Proceedings of the 20th International Conference on Autonomous Agents and MultiAgent Systems*, pp. 560–568, 2021.
- [30] R. Gorsane, O. Mahjoub, R. J. de Kock, R. Dubb, S. Singh, and A. Pretorius, “Towards a Standardised Performance Evaluation Protocol for Cooperative MARL,” in *Advances in Neural Information Processing Systems*, vol. 35, pp. 5510–5521, 2022.
- [31] A. Kapoor, B. Freed, J. Schneider, and H. Choset, “Assigning Credit with Partial Reward Decoupling in Multi-Agent Proximal Policy Optimization,” in *Proceedings of the Reinforcement Learning Conference*, 2024.
- [32] M. Akinmolayan, D. Josyula, and D. Casbeer, “Adaptive Interception in Dynamic Domains: Exploration of Hybrid Reinforcement Learning in Pursuit-Evasion Tasks,” *Proceedings of the AAAI Spring Symposium Series*, vol. 8, no. 1, pp. 2–10, 2026, doi: 10.1609/aaais.v8i1.42510.
- [33] X. Zhao and Y. Xie, “Multi-level Advantage Credit Assignment for Cooperative Multi-Agent Reinforcement Learning,” in *Proceedings of the 28th International Conference on Artificial Intelligence and Statistics*, PMLR, vol. 258, pp. 2971–2979, 2025.
- [34] B. Zhao, Q. Guo, X. Wang, R. Hedjam, and G. Zhong, “MA2MB: Multi-Agent Mutual-Advising Model-Based Reinforcement Learning for Pursuit and Evasion Games,” *Robotics and Autonomous Systems*, vol. 203, Art. no. 105530, 2026.